

ARTICLE TYPE

Chord Recognition in Symbolic Music: A Segmental CRF Model, Segment-Level Features, and Comparative Evaluations on Classical and Popular Music

Kristen Masada* and Razvan Bunescu*

Abstract

We present a new approach to harmonic analysis that is trained to segment music into a sequence of chord spans tagged with chord labels. Formulated as a semi-Markov Conditional Random Field (semi-CRF), this joint segmentation and labeling approach enables the use of a rich set of segment-level features, such as segment purity and chord coverage, that capture the extent to which the events in an entire segment of music are compatible with a candidate chord label. The new chord recognition model is evaluated extensively on three corpora of classical music and a newly created corpus of rock music. Experimental results show that the semi-CRF model performs substantially better than previous approaches when trained on a sufficient number of labeled examples and remains competitive when the amount of training data is limited.

Keywords: harmonic analysis, chord recognition, semi-CRF, segmental CRF, symbolic music.

1. Introduction and Motivation

Harmonic analysis is an important step towards creating high-level representations of tonal music. High-level structural relationships form an essential component of music analysis, whose aim is to achieve a deep understanding of how music works. At its most basic level, harmonic analysis of music in symbolic form requires the partitioning of a musical input into segments along the time dimension, such that the notes in each segment correspond to a musical chord. This *chord recognition* task can often be time consuming and cognitively demanding, hence the utility of computer-based implementations. Reflecting historical trends in artificial intelligence, automatic approaches to harmonic analysis have evolved from purely grammar-based and rule-based systems (Winoograd, 1968; Maxwell, 1992), to systems employing weighted rules and optimization algorithms (Temperley and Sleator, 1999; Pardo and Birmingham, 2002; Scholz and Ramalho, 2008; Rocher et al., 2009), to data driven approaches based on supervised machine learning (ML) (Raphael and Stoddard, 2003; Radicioni and Esposito, 2010). Due to their requirements for annotated data, ML approaches have also led to the development of music analysis datasets containing a large number of manually annotated harmonic structures, such as the 60 Bach chorales introduced in (Radicioni




Figure 1: Segment-based recognition (top) vs. event-based recognition (bottom) on measures 11 and 12 from Beethoven WoO68, using note onsets and offsets to create event boundaries.

and Esposito, 2010), and the 27 themes and variations of TAVERN (Devaney et al., 2015).

In this work, we consider the music to be in symbolic form, i.e. as a collection of notes specified in terms of onset, offset, pitch, and metrical position. Symbolic representations can be extracted from formats such as MIDI, kern, or MusicXML. A relatively common strategy in ML approaches to chord recognition in symbolic music is to break the musical input into a sequence of short duration spans and then train sequence tagging algorithms such as Hidden Markov Models (HMMs) to assign a chord label to each span in the sequence (bottom of Figure 1). The spans can result from quantization using a fixed musical period

*School of Electrical Engineering and Computer Science, Ohio University, Athens, OH

such as half a measure (Raphael and Stoddard, 2003). Alternatively, they can be constructed from consecutive note onsets and offsets (Radicioni and Esposito, 2010), as we also do in this paper. Variable-length chord segments are then created by joining consecutive spans labeled with the same chord symbol (at the top in Figure 1). A significant drawback of these short-span tagging approaches is that they do not explicitly model candidate segments during training and inference, consequently they cannot use segment-level features. Such features are needed, for example, to identify figuration notes (Appendix B) or to help label segments that do not start with the root note. The chordal analysis system of Pardo and Birmingham (2002) is an example where the assignment of chords to segments takes into account segment-based features, however the features have pre-defined weights and it uses a processing pipeline where segmentation is done independently of chord labeling.

In this paper, we propose a machine learning approach to chord recognition formulated under the framework of semi-Markov Conditional Random Fields (semi-CRFs). Also called segmental CRFs, this class of probabilistic graphical models can be trained to do joint segmentation and labeling of symbolic music (Section 2), using efficient Viterbi-based inference algorithms whose time complexity is linear in the length of the input. The system employs a set of chord labels (Section 3) that correspond to the main types of tonal music chords (Appendix A) found in the evaluation datasets. Compared to HMMs and sequential CRFs which label the events in a sequence, segmental CRFs label candidate segments, as such they can exploit segment-level features. Correspondingly, we define a rich set of features that capture the extent to which the events in an entire segment of music are compatible with a candidate chord label (Section 4). The semi-CRF model incorporating these features is evaluated on three classical music datasets and a newly created dataset of popular music (Section 5). Experimental comparisons with two previous chord recognition models show that segmental CRFs obtain substantial improvements in performance on the three larger datasets, while also being competitive with the previous approaches on the smaller dataset (Section 6).

2. Semi-CRF Model for Chord Recognition

Since harmonic changes may occur only when notes begin or end, we first create a sorted list of all the note onsets and offsets in the input music, i.e. the list of *partition points* (Pardo and Birmingham, 2002), shown as vertical dotted lines in Figure 1. A basic music *event* (Radicioni and Esposito, 2010) is then defined as the set of pitches sounding in the time interval between two consecutive partition points. As an example, Table 1 provides the pitches and overall duration for each event shown in Figure 2. The segment num-




Figure 2: Segment and labels (top) vs. events (bottom) for measure 12 from Beethoven WoO68.

Seg.	Label	Event	Pitches	Len.
s_1	G7	e_1	G3, B3, D4, G5	1/8
	G7	e_2	G3, B3, D4, F5	1/8
	G7	e_3	B4, D5	3/16
	G7	e_4	B4, D5	1/16
s_2	C	e_5	C4, C5, E5	1/8
	C	e_6	G3, C5, E5	1/8
	C	e_7	E3, G4, C5, E5	1/8
	C	e_8	C3, G4, C5, E5	1/8

Table 1: Input representation for measure 12 from Beethoven WoO68, showing the pitches and duration for each event, as well as the corresponding segment and label, where G7 stands for G:maj:add7, and C stands for C:maj.

ber and chord label associated with each event are also included. Not shown in this table is a boolean value for each pitch indicating whether or not it is held over from the previous event. For instance, this value would be false for C5 and E5 appearing in event e_5 , but true for C5 and E5 in event e_6 .

Let $\mathbf{s} = \langle s_1, s_2, \dots, s_K \rangle$ denote a segmentation of the musical input \mathbf{x} , where a segment $s_k = \langle s_k.f, s_k.l \rangle$ is identified by the positions $s_k.f$ and $s_k.l$ of its first and last events, respectively. Let $\mathbf{y} = \langle y_1, y_2, \dots, y_K \rangle$ be the vector of chord labels corresponding to the segmentation \mathbf{s} . A semi-Markov CRF (Sarawagi and Cohen, 2004) defines a probability distribution over segmentations and their labels as shown in Equations 1 and 2. Here, the global segmentation feature vector \mathbf{F} decomposes as a sum of local segment feature vectors $\mathbf{f}(s_k, y_k, y_{k-1}, \mathbf{x})$, with label y_0 set to a constant “no chord” value.

$$P(\mathbf{s}, \mathbf{y} | \mathbf{x}, \mathbf{w}) = \frac{e^{\mathbf{w}^T \mathbf{F}(\mathbf{s}, \mathbf{y}, \mathbf{x})}}{Z(\mathbf{x})} \quad (1)$$

$$\mathbf{F}(\mathbf{s}, \mathbf{y}, \mathbf{x}) = \sum_{k=1}^K \mathbf{f}(s_k, y_k, y_{k-1}, \mathbf{x}) \quad (2)$$

where $Z(\mathbf{x}) = \sum_{\mathbf{s}', \mathbf{y}'} e^{\mathbf{w}^T \mathbf{F}(\mathbf{s}', \mathbf{y}', \mathbf{x})}$ and \mathbf{w} is a vector of parameters.

Following Muis and Lu (Muis and Lu, 2016), for faster inference, we further restrict the local segment

features to two types: *segment-label features* $\mathbf{f}(s_k, y_k, \mathbf{x})$ that depend on the segment and its label, and label *transition features* $\mathbf{g}(y_k, y_{k-1}, \mathbf{x})$ that depend on the labels of the current and previous segments. The corresponding probability distribution over segmentations is shown in Equations 3 to 5, which use two vectors of parameters: \mathbf{w} for segment-label features and \mathbf{u} for transition features.

$$P(\mathbf{s}, \mathbf{y} | \mathbf{x}, \mathbf{w}, \mathbf{u}) = \frac{e^{\mathbf{w}^T \mathbf{F}(\mathbf{s}, \mathbf{y}, \mathbf{x}) + \mathbf{u}^T \mathbf{G}(\mathbf{s}, \mathbf{y}, \mathbf{x})}}{Z(\mathbf{x})} \quad (3)$$

$$\mathbf{F}(\mathbf{s}, \mathbf{y}, \mathbf{x}) = \sum_{k=1}^K \mathbf{f}(s_k, y_k, \mathbf{x}) \quad (4)$$

$$\mathbf{G}(\mathbf{s}, \mathbf{y}, \mathbf{x}) = \sum_{k=1}^K \mathbf{g}(y_k, y_{k-1}, \mathbf{x}) \quad (5)$$

Given an arbitrary segment s and a label y , the vector of segment-label features can be written as $\mathbf{f}(s, y, \mathbf{x}) = [f_1(s, y), \dots, f_{|\mathbf{f}|}(s, y)]$, where the input \mathbf{x} is left implicit in order to compress the notation. Similarly, given arbitrary labels y and y' , the vector of label transition features can be written as $\mathbf{g}(y, y', \mathbf{x}) = [g_1(y, y'), \dots, g_{|\mathbf{g}|}(y, y')]$. In Section 4 we describe the set of segment-label features $f_i(s, y)$ and label transition features $g_j(y, y')$ that are used in our semi-CRF chord recognition system.

As probabilistic graphical models, semi-CRFs can be represented using factor graphs, as illustrated in Figure 3. Factor graphs (Kschischang et al., 2001) are bipartite graphs that express how a global function (e.g. $P(\mathbf{s}, \mathbf{y} | \mathbf{x}, \mathbf{w}, \mathbf{u})$) of many variables (e.g. s_k, y_k , and \mathbf{x}) factorizes into a product of local functions, or *factors*, (e.g. f and g) defined over fewer variables.




Figure 3: Factor graph representation of the semi-CRF.

Equations 4 and 5 show that the contribution of any given feature to the final log-likelihood score is given by summing up its value over all the segments (for local features f) or segment pairs (for local features g). This design choice stems from two assumptions. First, we adopt the stationarity assumption, according to which the segment-label feature distribution does not change with the position in the music. Second, we use the Markov assumption, which implies that the label of a segment depends only on its boundaries and the labels of the adjacent segments. This assumption leads to the factorization of the probability distribution into a product of potentials. Both the stationarity assumption and the Markov assumption are

commonly used in ML models for structured outputs, such as linear CRFs (Lafferty et al., 2001), semi-CRFs (Sarawagi and Cohen, 2004), HMMs (Rabiner, 1989), structural SVMs (Tsochantaridis et al., 2004), or the structured perceptron (Collins, 2002) used in HMPerceptron. These assumptions lead to summing the same feature over multiple substructures in the overall output score, which makes *inference* and *learning* tractable using dynamic programming.

The *inference* problem for semi-CRFs refers to finding the most likely segmentation $\hat{\mathbf{s}}$ and its labeling $\hat{\mathbf{y}}$ for an input \mathbf{x} , given the model parameters. For the weak semi-CRF model in Equation 3, this corresponds to:

$$\hat{\mathbf{s}}, \hat{\mathbf{y}} = \arg \max_{\mathbf{s}, \mathbf{y}} P(\mathbf{s}, \mathbf{y} | \mathbf{x}, \mathbf{w}, \mathbf{u}) \quad (6)$$

$$= \arg \max_{\mathbf{s}, \mathbf{y}} \mathbf{w}^T \mathbf{F}(\mathbf{s}, \mathbf{y}, \mathbf{x}) + \mathbf{u}^T \mathbf{G}(\mathbf{s}, \mathbf{y}, \mathbf{x}) \quad (7)$$

$$= \arg \max_{\mathbf{s}, \mathbf{y}} \mathbf{w}^T \sum_{k=1}^K \mathbf{f}(s_k, y_k, \mathbf{x}) + \mathbf{u}^T \sum_{k=1}^K \mathbf{g}(y_k, y_{k-1}, \mathbf{x}) \quad (8)$$

The maximum is taken over all possible labeled segmentations of the input, up to a maximum segment length. Correspondingly, \mathbf{s} and \mathbf{y} can be seen as “candidate” segmentations and “candidate” labelings, respectively. Their number is exponential in the length of the input, which rules out a brute-force search. However, due to the factorization into vectors of local features $f_i(s, y)$ and $g_j(y, y')$, it can be shown that the optimization problem from Equation 8 can be solved with a semi-Markov analogue of the usual Viterbi algorithm. Let L be a maximum segment length. Following (Sarawagi and Cohen, 2004), let $V(i, y)$ denote the largest value $\mathbf{w}^T \mathbf{F}(\tilde{\mathbf{s}}, \tilde{\mathbf{y}}, \mathbf{x}) + \mathbf{u}^T \mathbf{G}(\tilde{\mathbf{s}}, \tilde{\mathbf{y}}, \mathbf{x})$ of a partial segmentation $\tilde{\mathbf{s}}$ such that its last segment ends at position i and has label y . Then $V(i, y)$ can be computed with the following dynamic programming recursion for $i = 1, 2, \dots, |\mathbf{x}|$:

$$V(i, y) = \max_{y', 1 \leq l \leq L} V(i-l, y') + \mathbf{w}^T \mathbf{f}(\langle i-l+1, i \rangle, y, \mathbf{x}) + \mathbf{u}^T \mathbf{g}(y, y', \mathbf{x}) \quad (9)$$

where the base cases are $V(0, y) = 0$ and $V(j, y) = -\infty$ if $j < 0$, and $\langle i-l+1, i \rangle$ denotes the segment starting at position $i-l+1$ and ending at position i . Once $V(|\mathbf{x}|, y)$ is computed for all labels y , the best labeled segmentation can be recovered in linear time by following the path traced by $\max_y V(|\mathbf{x}|, y)$.

The *learning* problem for semi-CRFs refers to finding the model parameters that maximize the likelihood over a set of training sequences $T = \{\mathbf{x}_n, \mathbf{s}_n, \mathbf{y}_n\}_{n=1}^N$. Usually this is done by minimizing the negative log-likelihood $-L(T; \mathbf{w}, \mathbf{u})$ and an L2 regularization term, as shown below for weak semi-CRFs:

$$L(T; \mathbf{w}, \mathbf{u}) = \sum_{n=1}^N \mathbf{w}^T \mathbf{F}(\mathbf{s}_n, \mathbf{y}_n, \mathbf{x}_n) + \mathbf{u}^T \mathbf{G}(\mathbf{s}_n, \mathbf{y}_n, \mathbf{x}_n) - \log Z(\mathbf{x}_n) \quad (10)$$

$$\hat{\mathbf{w}}, \hat{\mathbf{u}} = \arg \min_{\mathbf{w}, \mathbf{u}} -L(T; \mathbf{w}, \mathbf{u}) + \frac{\lambda}{2} (\|\mathbf{w}\|^2 + \|\mathbf{u}\|^2) \quad (11)$$

This is a convex optimization problem, which is solved with the L-BFGS procedure in the StatNLP package used to implement our system. The partition function $Z(\mathbf{x})$ and the feature expectations that appear in the gradient of the objective function are computed efficiently using a dynamic programming algorithm similar to the forward-backward procedure (Sarawagi and Cohen, 2004).

3. Chord Recognition Labels

A *chord* is a group of notes that form a cohesive harmonic unit to the listener when sounding simultaneously (Aldwell et al., 2011). As explained in Appendix A, we design our system to handle the following types of chords: *triads*, *augmented 6th chords*, *suspended chords*, and *power chords*. The chord labels used in previous chord recognition research range from coarse grained labels that indicate only the chord root (Temperley and Sleator, 1999) to fine grained labels that capture mode, inversions, added and missing notes (Harte, 2010), and even chord function (Devaney et al., 2015). Here we follow the middle ground proposed by Radicioni and Esposito (2010) and define a core set of labels for triads that encode the chord root (12 pitch classes), the mode (major, minor, diminished), and the added note (none, fourth, sixth, seventh), for a total of 144 different labels. For example, the label *C-major-none* for a simple C major triad corresponds to the combination of a root of C with a mode of *major* and no added note. This is different from the label *C-major-seventh* for a C major seventh chord, which corresponds to the combination of a root of C with a mode of *major* and an added note of *seventh*. Note that there is only one generic type of added seventh note, irrespective of whether the interval is a major, minor, or diminished seventh, which means that a C major seventh chord and a C dominant seventh chord are mapped to the same label. However, once the system recognizes a chord with an added seventh, determining whether it is a major, minor, or diminished seventh can be done accurately in a simple post-processing step: determine if the chord contains a non figuration note (defined in Appendix B) that is 11, 10, or 9 half steps from the root, respectively, inverted or not, modulo 12. Once the type of the seventh interval is determined, it is straightforward to determine the type of seventh chord (dominant, major, minor, minor-major, fully diminished, or half-diminished) based on the mode of the chord (major, minor, or diminished).

Augmented sixth chords are modeled through a set of 36 labels that capture the lowest note (12 pitch classes) and the 3 types (Appendix A.2). Similarly, suspended and power chords are modeled through a set of 48 labels that capture the root note (12 pitch classes) and the 4 types (Appendix A.3).

Because the labels do not encode for function, the model does not require knowing the key in which the

input was written. While the number of labels may seem large, the number of parameters in our model is largely independent of the number of labels. This is because we design the chord recognition features (Section 4) to not test for the chord root, which also enables the system to recognize chords that were not seen during training. The decision to not use the key context was partly motivated by the fact that 3 of the 4 datasets we used for experimental evaluation do not have functional annotations (see Section 5). Additionally, complete key annotation can be difficult to perform, both manually and automatically. Key changes occur gradually, thus making it difficult to determine the exact location where one key ends and another begins (Papadopoulos and Peeters, 2009). This makes locating modulations and tonicizations difficult and also hard to evaluate (Gómez, 2006). At the same time, we recognize that harmonic analysis is not complete without functional analysis. Functional analysis features could also benefit the basic chord recognition task described in this paper. In particular, the chord transition features that we define in Appendix C.4 depend on the absolute distance in half steps between the roots of the chords. However, a V-I transition has a different distribution than a I-IV transition, even though the root distance is the same. Chord transition distributions also differ between minor and major keys. As such, using key context could further improve chord recognition.

4. Chord Recognition Features

The semi-CRF model uses five major types of features, as described in detail in Appendix C. Segment purity features compute the percentage of segment notes that belong to a given chord (Appendix C.1). We include these on the grounds that segments with a higher purity with respect to a chord are more likely to be labeled with that chord. Chord coverage features determine if each note in a given chord appears at least once in the segment (Appendix C.2). Similar to segment purity, if the segment covers a higher percentage of the chord’s notes, it is more likely to be labeled with that chord. Bass features determine which note of a given chord appears as the bass in the segment (Appendix C.3). For a correctly labeled segment, its bass note often matches the root of its chord label. If the bass note instead matches the chord’s third or fifth, or is an added dissonance, this may indicate that the chord y is inverted or incorrect. Chord bigram features capture chord transition information (Appendix C.4). These features are useful in that the arrangement of chords in chord progressions is an important component of *harmonic syntax*. Finally, we include metrical accent features for chord changes, as chord segments are more likely to begin on accented beats (Appendix C.5).

5. Chord Recognition Datasets

For evaluation, we used four chord recognition datasets:

1. BaCh: this is the Bach Choral Harmony Dataset, a corpus of 60 four-part Bach chorales that contains 5,664 events and 3,090 segments in total (Radicioni and Esposito, 2010).
2. TAVERN: this is a corpus of 27 complete sets of themes and variations for piano, composed by Mozart and Beethoven. It consists of 63,876 events and 12,802 segments overall (Devaney et al., 2015).
3. KP Corpus: the Kostka-Payne corpus is a dataset of 46 excerpts compiled by Bryan Pardo from Kostka and Payne’s music theory textbook. It contains 3,888 events and 911 segments (Kostka and Payne, 1984).
4. Rock: this is a corpus of 59 pop and rock songs that we compiled from Hal Leonard’s *The Best Rock Songs Ever (Easy Piano)* songbook. It is 25,621 events and 4,221 segments in length.

5.1 The Bach Chorale (BaCh) Dataset

The BaCh corpus has been annotated by a human expert with chord labels, using the set of triad labels described in Section 3. Of the 144 possible labels, 102 appear in the dataset and of these only 68 appear 5 times or more. Some of the chord labels used in the manual annotation are enharmonic, e.g. C-sharp major and D-flat major, or D-sharp major and E-flat major. Reliably producing one of two enharmonic chords cannot be expected from a system that is agnostic of the key context. Therefore, we normalize the chord labels and for each mode we define a set of 12 canonical roots, one for each scale degree. When two enharmonic chords are available for a given scale degree, we selected the one with the fewest sharps or flats in the corresponding key signature. Consequently, for the major mode we use the canonical root set {C, Db, D, Eb, E, F, Gb, G, Ab, A, Bb, B}, whereas for the minor and diminished modes we used the root set {C, C#, D, D#, E, F, F#, G, G#, A, Bb, B}. Thus, if a chord is manually labeled as C-sharp major, the label is automatically changed to the enharmonic D-flat major. The actual chord notes used in the music are left unchanged. Whether they are spelled with sharps or flats is immaterial, as long as they are enharmonic with the root, third, fifth, or added note of the labeled chord. After performing enharmonic normalization on the chords in the dataset, 90 labels remain.

5.2 The TAVERN Dataset

The TAVERN dataset¹ currently contains 17 works by Beethoven (181 variations) and 10 by Mozart (100 variations). The themes and variations are divided into a total of 1,060 phrases, 939 in major and 121 in minor. The pieces have two levels of segmentations: chords and phrases. The chords are annotated with

Roman numerals, using the Humdrum representation for functional harmony². When finished, each phrase will have annotations from two different experts, with a third expert adjudicating cases of disagreement between the two. After adjudication, a unique annotation of each phrase is created and joined with the note data into a combined file encoded in standard `**kern` format. However, many pieces do not currently have the second annotation or the adjudicated version. Consequently, we only used the first annotation for each of the 27 sets. Furthermore, since our chord recognition approach is key agnostic, we developed a script that automatically translated the Roman numeral notation into the key-independent canonical set of labels used in BaCh. Because the TAVERN annotation does not mark added fourth notes, the only added chords that were generated by the translation script were those containing sixths and sevenths. This results in a set of 108 possible labels, of which 69 appear in the dataset.

5.3 The Kostka and Payne Corpus

The Kostka-Payne (KP) corpus³ does not contain chords with added fourth or sixth notes. However, it includes fine-grained chord types that are outside of the label set of triads described in Section 3, such as fully and half-diminished seventh chords, dominant seventh chords, and dominant seventh flat ninth chords. We map these seventh chord variants to the generic added seventh chords, as discussed in Section 3. Chords with ninth intervals are mapped to the corresponding chord without the ninth in our label set. The KP Corpus also contains the three types of augmented 6th chords introduced in Appendix A.2. Thus, by extending our chord set to include augmented 6th labels, there are $12 \text{ roots} \times 3 \text{ triad modes} \times 2 \text{ added notes} + 12 \text{ bass notes} \times 3 \text{ aug6 modes} = 108$ possible labels overall. Of these, 76 appear in the dataset.

A number of MIDI files in the KP corpus contain unlabeled sections at the beginning of the song. These sections also appear as unlabeled in the original Kostka-Payne textbook. We omitted these sections from our evaluation, and also did not include them in the KP Corpus event and segment counts. Bryan Pardo’s original MIDI files for the KP Corpus also contain several missing chords, as well as chord labels that are shifted from their true onsets. We used chord and beat list files sent to us by David Temperley to correct these mistakes.

5.4 The Rock Dataset

To evaluate the system’s ability to recognize chords in a different genre, we compiled a corpus of 59 pop and rock songs from Hal Leonard’s *The Best Rock Songs Ever (Easy Piano)* songbook. Like the KP Corpus, the Rock dataset contains chords with added ninths—including major ninth chords and dominant seventh chords with a sharpened ninth—as well as inverted chords. We omit the ninth and inversion numbers in these cases.

Unique from the other datasets, the Rock dataset also possesses suspended and power chords. We extend our chord set to include these, adding suspended second, suspended fourth, dominant seventh suspended fourth, and power chords. We use the major mode canonical root set for suspended second and power chords and the minor canonical root set for suspended fourth chords, as this configuration produces the least number of accidentals. In all, there are $12 \text{ roots} \times 3 \text{ triad modes} \times 4 \text{ added notes} + 12 \text{ roots} \times 4 \text{ sus and pow modes} = 192$ possible labels, with only 48 appearing in the dataset.

Similar to the KP Corpus, unlabeled segments occur at the beginning of some songs, which we omit from evaluation. Additionally, the Rock dataset uses an N.C. (i.e. no chord) label for some segments within songs where the chord is unclear. We broke songs containing this label into subsections consisting of the segments occurring before and after each N.C. segment, discarding subsections less than three measures long.

To create the Rock dataset, we converted printed sheet music to MusicXML files using the optical music recognition (OMR) software PhotoScore⁴. We noticed in the process of making the dataset that some of the originally annotated labels were incorrect. For instance, some segments with added note labels were missing the added note, while other segments were missing the root or were labeled with an incorrect mode. We automatically detected these cases and corrected each label by hand, considering context and genre-specific theory. We also omitted two songs (“Takin’ Care of Business” and “I Love Rock N’ Roll”) from the 61 songs in the original Hal Leonard songbook, the former because of its atonality and the latter because of a high percentage of mistakes in the original labels.

6. Experimental Evaluation

We implemented the semi-Markov CRF chord recognition system using a multi-threaded package⁵ that has been previously used for noun-phrase chunking of informal text (Muis and Lu, 2016). The following sections describe the experimental results obtained on the four datasets from Section 5 for: our semi-CRF system; Radicioni and Esposito’s perceptron-trained HMM system, HMPerceptron; and Temperley’s computational music system, Melisma Music Analyzer⁶. When interpreting these results, it is important to consider a number of important differences among the three systems:

- HMPerceptron and semi-CRF are data driven, therefore their performance depends on the number of training examples available. Both approaches are agnostic of music theoretic principles such as harmony changing primarily on strong metric positions, however they can learn such tendencies to the extent they are present in the training data.

- Compared to HMPerceptron, semi-CRFs can use segment-level features. Besides this conceptual difference, the semi-CRF system described here uses a much larger number of features than the HMPerceptron system, which by itself can lead to better performance but may also require more training examples.
- Both Melisma and HMPerceptron use metrical accents automatically induced by Melisma, whereas semi-CRF uses the Music21 accents derived from the notated meter. The more accurate notated meter could favor the semi-CRF system, although results in Section 6.1 show that, at least on BaCh, HMPerceptron does not benefit from using the notated meter.

Table 2 shows a summary of the full chord and root-level experimental results provided in this section. Two overall types of measures are used to evaluate a system’s performance on a dataset: event-level accuracy (Acc_E) and segment-level F-measure (F_S). Acc_E simply refers to the percentage of events for which the system predicts the correct label out of the total number of events in the dataset. Segment-level F-measure is computed based on precision and recall, two evaluation measures commonly used in information retrieval (Baeza-Yates and Ribeiro-Neto, 1999), as follows:

- Precision (P_S) is the percentage of segments predicted correctly by the system out of the total number of segments that it predicts (correctly or incorrectly) for all songs in the dataset.
- Recall (R_S) is the percentage of segments predicted correctly out of the total number of segments annotated in the original score for all songs in the dataset.
- F-Measure (F_S) is the harmonic mean between P_S and R_S , i.e. $F_S = 2P_S R_S / (P_S + R_S)$.

Note that a predicted segment is considered correct if and only if both its boundaries and its label match those of a true segment.

6.1 BaCh Evaluation

We evaluated the semi-CRF model on BaCh using 10-fold cross validation: the 60 Bach chorales were randomly split into a set of 10 folds, and each fold was used as test data, with the other nine folds being used for training. We then evaluated HMPerceptron using the same randomly generated folds to enable comparison with our system. However, we noticed that the performance of HMPerceptron could vary significantly between two different random partitions of the data into folds. Therefore, we repeated the 10-fold cross validation experiment 10 times, each time shuffling the 60 Bach chorales and partitioning them into 10 folds. For each experiment, the test results from the 10 folds were pooled together and one value was computed for each performance measure (accuracy, precision, recall, and F-measure). The overall performance measures for

	Statistics			Full chord evaluation				Root-level evaluation					
				semi-CRF		HMPerceptron		semi-CRF		HMPerceptron		Melisma	
Dataset	Events	Seg.'s	Labels	Acc _E	F _S	Acc _E	F _S	Acc _E	F _S	Acc _E	F _S	Acc _E	F _S
BaCh	5,664	3,090	90	83.2	77.5	77.2	69.9	88.9	84.2	84.8	77.0	84.3	74.7
TAVERN	63,876	12,802	69	78.0	64.0	57.0	22.5	86.0	71.4	69.2	33.2	76.7	41.5
KPCorpus	3,888	911	76	73.0	53.0	72.9	45.4	79.3	59.0	79.0	51.9	81.9	62.2
Rock	25,621	4,221	48	70.1	55.9	61.3	34.6	86.1	65.1	80.7	42.9	77.9	36.3

Table 2: Dataset statistics and summary of results (event-level accuracy Acc_E and segment-level F-measure F_S).

the two systems were then computed by averaging over the 10 values (one from each experiment). The sample standard deviation for each performance measure was also computed over the same 10 values.

For semi-CRF, we computed the frequency of occurrence of each feature in the training data, using only the true segment boundaries and their labels. To speedup training and reduce overfitting, we only used features whose counts were at least 5. The performance measures were computed by averaging the results from the 10 test folds for each of the fold sets. Table 3 shows the averaged event-level and segment-level performance of the semi-CRF model, together with two versions of the HMPerceptron: HMPerceptron₁, for which we do enharmonic normalization both on training and test data, similar to the normalization done for semi-CRF; and HMPerceptron₂, which is the original system from (Radicioni and Esposito, 2010) that does enharmonic normalization only on test data.

BaCh: Full chord evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	83.2 0.2	79.4 0.2	75.8 0.2	77.5 0.2
HMPerceptron ₁	77.2 2.1	71.2 2.0	68.8 2.2	69.9 1.8
HMPerceptron ₂	77.0 2.1	71.0 2.0	68.5 2.3	69.7 1.8

Table 3: Comparative results (%) and standard deviations on the BaCh dataset, using Event-level **accuracy** (Acc_E) and Segment-level **precision** (P_S), **recall** (R_S), and **F-measure** (F_S).

The semi-CRF model achieves a 6.2% improvement in event-level accuracy over the original model HMPerceptron₂, which corresponds to a 27.0% relative error reduction¹. The improvement in accuracy over HMPerceptron₁ is statistically significant at an averaged p -value of 0.001, using a one-tailed Welch’s t -test over the sample of 60 chorale results for each of the 10 fold sets. The improvement in segment-level performance is even more substantial, with a 7.8% absolute improvement in F-measure over the original HMPerceptron₂ model, and a 7.6% improvement in F-measure over the HMPerceptron₁ version, which is statistically significant at an averaged p -value of

0.002, using a one-tailed Welch’s t -test. The standard deviation values computed for both event-level accuracy and F-Measure are about one order of magnitude smaller for semi-CRF than for HMPerceptron, demonstrating that the semi-CRF is also more stable than the HMPerceptron. As HMPerceptron₁ outperforms HMPerceptron₂ in both event and segment-level accuracies, we will use HMPerceptron₁ for the remaining evaluations and will simply refer to it as HMPerceptron.

BaCh: Root only evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	88.9	85.4	83.0	84.2
HMPerceptron	84.8	78.0	76.2	77.0
Melisma	84.3	73.2	76.3	74.7

Table 4: Root only results (%) on the BaCh dataset, using Event-level **accuracy** (Acc_E) and Segment-level **precision** (P_S), **recall** (R_S), and **F-measure** (F_S).

We also evaluated performance in terms of predicting the correct root of the chord, e.g. if the true chord label were C:maj, a predicted chord of C:maj:add7 would still be considered correct, because it has the same root as the correct label. We performed this evaluation for semi-CRF, HMPerceptron, and the harmony component of Temperley’s Melisma. Results show that semi-CRF improves upon the event-level accuracy of HMPerceptron by 4.1%, producing a relative error reduction of 27.0%, and that of Melisma by 4.6%. Semi-CRF also achieves an F-measure that is 7.2% higher than HMPerceptron and 9.5% higher than Melisma. These improvements are statistically significant with a p -value of 0.01 using a one-tailed Welch’s t -test.

BaCh: Metrical accent evaluation of semi-CRF				
System	Acc _E	P _S	R _S	F _S
With accent	83.6	79.6	75.9	77.6
Without accent	77.7	74.8	68.0	71.2

Table 5: Full chord Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the BaCh dataset, with and without metrical accent features.

Metrical accent is important for harmonic analysis: chord changes tend to happen in strong metrical positions; figuration such as passing and neighboring tones

¹27% = (83.2 – 77.0)/(100 – 77.0)

appear in metrically weak positions, whereas suspensions appear on metrically strong beats. We verified empirically the importance of metrical accent by evaluating the semi-CRF model on a random fold set from the BaCh corpus with and without all accent-based features. The results from Table 5 show a substantial decrease in accuracy when the accent-based features are removed from the system.

Finally, we ran an evaluation of HMPerceptron on a random fold set from BaCh in two scenarios: HMPerceptron with Melisma metrical accent and HMPerceptron with Music21 accent. The results did not show a significant difference: with Melisma accent the event accuracy was 79.8% for an F-measure of 70.2%, whereas with Music21 accent the event accuracy was 79.8% for an F-measure of 70.3%. This negligible difference is likely due to the fact that HMPerceptron uses only coarse-grained accent information, i.e. whether a position is accented (Melisma accent 3 or more) or not accented (Melisma accent less than 3).

6.1.1 BaCh Error Analysis

Error analysis revealed wrong predictions being made on chords that contained dissonances that spanned the duration of the entire segment (e.g. a second above the root of the annotated chord), likely due to an insufficient number of such examples during training. Manual inspection also revealed a non-trivial number of cases in which we disagreed with the manually annotated chords, e.g. some chord labels were clear mistakes, as they did not contain any of the notes in the chord. This further illustrates the necessity of building music analysis datasets that are annotated by multiple experts, with adjudication steps akin to the ones followed by TAVERN.

6.2 TAVERN Evaluation

To evaluate on the TAVERN corpus, we created a fixed training-test split: 6 Beethoven sets (*B063*, *B064*, *B065*, *B066*, *B068*, *B069*) and 4 Mozart sets (*K025*, *K179*, *K265*, *K353*) were used for testing, while the remaining 11 Beethoven sets and 6 Mozart sets were used for training. All sets were normalized enharmonically before being used for training or testing. Table 6 shows the event-level and segment-level performance of the semi-CRF and HMPerceptron model on the TAVERN dataset.

TAVERN: Full chord evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	78.0	67.3	60.9	64.0
HMPerceptron	57.0	24.5	20.8	22.5

Table 6: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the TAVERN dataset.

As shown in Table 6, semi-CRF outperforms HMPerceptron by 21.0% for event-level chord evaluation and




Figure 4: Semi-CRF correctly predicts A:maj7 (top) for the first beat of measure 55 from Mozart K025, while HMPerceptron predicts C#:dim (bottom).




Figure 5: Semi-CRF correctly predicts C:maj (top) for all of measure 280 from Mozart K179, while HMPerceptron predicts E:min (bottom) for the first beat and C:maj for the other two beats (bottom).

by 41.5% in terms of chord-level F-measure. Root only evaluations provided in Table 7 reveal that semi-CRF improves upon HMPerceptron’s event-level root accuracy by 16.8% and Melisma’s event accuracy by 9.3%. Semi-CRF also produces a segment-level F-measure value that is 38.2% higher than that of HMPerceptron and 29.9% higher than that of Melisma. These improvements are statistically significant with a *p*-value of 0.01 using a one-tailed Welch’s t-test.

TAVERN: Root only evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	86.0	74.6	68.4	71.4
HMPerceptron	69.2	38.2	29.4	33.2
Melisma	76.7	42.3	40.7	41.5

Table 7: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the TAVERN dataset.

6.2.1 TAVERN Error Analysis

The results in Tables 3 and 6 show that chord recognition is substantially more difficult in the TAVERN dataset than in BaCh. The comparatively lower performance on TAVERN is likely due to the substantially larger number of figurations and higher rhythmic diversity of the variations compared to the easier, mostly note-for-note texture of the chorales. Error analysis on TAVERN revealed many segments where the first event did not contain the root of the chord, such as in Figures 4 and 5. For such segments, HMPerceptron

incorrectly assigned chord labels whose root matched the bass of this first event. Since a single wrongly labeled event invalidates the entire segment, this can explain the larger discrepancy between the event-level accuracy and the segment-level performance. In contrast, semi-CRF assigned the correct labels in these cases, likely due to its ability to exploit context through segment-level features, such as the chord root coverage feature f_4 and its duration-weighted version f_{11} . In the case of Figure 4, C# appears in the bass of the first beat of the measure and HMPeception incorrectly predicts a segment with label C#:dim for this beat. In contrast, semi-CRF correctly predicts the label A:maj7 for this segment. In Figure 5, semi-CRF correctly predicts a C:maj segment that lasts for the entirety of the measure, while HMPeception predicts an E:min segment for the first beat, as E appears doubled in the bass here.

6.3 KP Corpus Evaluation

To evaluate on the full KP Corpus dataset, we split the songs into 11 folds. In this configuration, 9 folds contain 4 songs each, while the remaining 2 folds contain 5 songs. We then created two versions of semi-CRF: the original system without augmented 6th chord features (semi-CRF₁) and a system with augmented 6th features (semi-CRF₂). We tested both versions on all 46 songs, as shown in Table 8. We could not perform the same evaluation on HMPeception because it was not designed to handle augmented 6th chords.

KP Corpus 46 songs: Full chord evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF ₁	72.0	59.0	49.2	53.5
semi-CRF ₂	73.4	59.6	50.1	54.3

Table 8: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the KP Corpus dataset.

The results in Table 8 demonstrate the utility of adding augmented 6th chord features to our system, as semi-CRF₂ outperforms semi-CRF₁ on all measures. We will use semi-CRF₂ for the rest of the evaluations in this section, simply calling it semi-CRF.

We additionally perform root only evaluation on the full dataset for semi-CRF and Melisma. We ignore events that belong to the true augmented 6th chord segments when computing the root accuracies for both systems, as augmented 6th chords technically do not contain a root note. As shown in Table 9, Melisma is only marginally better than semi-CRF in terms of event-level root accuracy, however it has a segment-level F-measure that is 1.1% better.

To enable comparison with HMPeception, we also evaluate all systems on the 36 songs that do not contain augmented 6th chords. Because of the reduced number of songs available for training, we used leave-one-out evaluation for both semi-CRF and HMPeception.

KP Corpus 46 songs: Root only evaluation

System	Acc _E	P _S	R _S	F _S
semi-CRF	80.7	66.3	56.2	60.8
Melisma	80.9	60.6	63.3	61.9

Table 9: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the KP Corpus dataset.

tron. Table 10 shows that semi-CRF obtains a marginal improvement in chord event accuracy and a more substantial 7.6% improvement in segment-level F-measure in comparison with HMPeception. The comparative results in Table 11 show that Melisma outperforms both machine learning systems for root only evaluation. Nevertheless, the semi-CRF is still competitive with Melisma in terms of both event-level accuracy and segment-level F-measure.

KP Corpus 36 songs: Full chord evaluation

System	Acc _E	P _S	R _S	F _S
semi-CRF	73.0	55.6	50.7	53.0
HMPeception	72.9	48.2	43.6	45.4

Table 10: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the KP Corpus dataset.

KP Corpus 36 songs: Root only evaluation

System	Acc _E	P _S	R _S	F _S
semi-CRF	79.3	61.8	56.4	59.0
HMPeception	79.0	54.7	49.9	51.9
Melisma	81.9	60.7	63.7	62.2

Table 11: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the KP Corpus dataset.

We additionally compare semi-CRF against the HarmAn algorithm created by Pardo and Birmingham (2002), which achieves a 75.8% event-level accuracy on the KP Corpus. We made several modifications to the initial evaluation of semi-CRF on the full KP Corpus to enable this comparison. For instance, Pardo and Birmingham omit a Schumann piece from their evaluation, testing HarmAn on 45 songs instead of 46. We omitted this piece as well. They also look at the labels that appear in the dataset beforehand, ignoring any segments whose correct labels are chords that appear less than 2% of the time when rounded. We followed suit with this, ignoring segments labeled with augmented 6th chords and other less common labels. Overall, semi-CRF obtains an event-level accuracy of 75.3%, demonstrating that it is competitive with HarmAn. However, it is important to note that these results are still not fully comparable: sometimes HarmAn predicts multiple labels for a single segment, and when the correct label is among these, Pardo and Birmingham divide by the number of labels the system predicts and

consider this fractional value to be correct. In contrast, semi-CRF always predicts one label per segment.

6.3.1 KP Corpus Error Analysis

Both machine learning systems struggled on the KP corpus, with Melisma performing better on both event-level accuracy and segment-level F-measure. This can be explained by the smaller dataset, and thus the smaller number of available training examples. The KP corpus was the smallest of the four datasets, especially in terms of the number of segments – less than a third compared to BaCh, and less than a tenth compared to TAVERN. Furthermore, the textbook excerpts are more diverse, as they are taken from 11 composers and are meant to illustrate a wide variety of music theory concepts, leading to mismatch between the training and test distributions and thus lower test performance.

6.4 Rock Evaluation

We split the 59 songs in the rock dataset into 10 folds: 9 folds with 6 songs and 1 fold with 5 songs. Similar to the full KP Corpus evaluation from Section 6.3, we create two versions of the semi-CRF model. The first is the original semi-CRF system (semi-CRF₁) which does not contain suspended and power chord features. The second is a new version of semi-CRF (semi-CRF₃) which has suspended and power chord features added to it. We do not include HMPerceptron in the evaluation of the full dataset, as it is not designed for suspended and power chords.

Rock 59 songs: Full chord evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF ₁	66.0	49.8	47.3	48.5
semi-CRF ₃	69.4	62.0	54.9	58.3

Table 12: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the Rock dataset.

As shown in Table 12, semi-CRF₃ obtains higher event and segment-level accuracies than semi-CRF₁. Therefore, we use semi-CRF₃ for the rest of the experiments, simply calling it semi-CRF.

We perform root only evaluation on the full Rock dataset using semi-CRF and Melisma. In this case, it is not necessary to omit the true segments whose labels are suspended or power chords, as these types of chords contain a root. As shown in Table 13, semi-CRF outperforms Melisma on all measures: it obtains a 8.4% improvement in event-level root accuracy and a 31.5% improvement in segment-level F-measure over Melisma.

We also evaluate only on the 51 songs that do not contain suspended or power chords to compare semi-CRF against HMPerceptron. We do this by splitting the reduced number of songs into 10 folds: 9 folds with 5 test songs and 46 training songs, and 1 fold with 6 test songs and 45 training songs. The results shown in

Rock 59 songs: Root only evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	85.8	70.9	63.2	66.8
Melisma	77.4	29.5	44.0	35.3

Table 13: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the Rock dataset.

Table 14 demonstrate that semi-CRF performs better than HMPerceptron: it achieves an 8.8% improvement in event-level chord accuracy and a 21.3% improvement in F-measure over HMPerceptron. Additionally, we evaluate the root-level performance of all systems on the 51 songs. The results in Table 15 show that the semi-CRF achieves better root-level accuracy than both systems: it obtains a 5.4% improvement in event-level root accuracy over HMPerceptron and a 8.2% improvement over Melisma. In terms of segment-level accuracy, it demonstrates a 22.2% improvement in F-measure over HMPerceptron and a 28.8% improvement over Melisma. These results are statistically significant with a p -value of 0.01 using a one-tailed Welch’s t-test.

Rock 51 songs: Full chord evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	70.1	58.8	53.2	55.9
HMPerceptron	61.3	41.0	29.9	34.6

Table 14: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the Rock dataset.

Rock 51 songs: Root only evaluation				
System	Acc _E	P _S	R _S	F _S
semi-CRF	86.1	68.6	61.9	65.1
HMPerceptron	80.7	51.3	36.9	42.9
Melisma	77.9	30.6	45.8	36.3

Table 15: Event (Acc_E) and Segment-level (P_S, R_S, F_S) results (%) on the Rock dataset.

6.4.1 Rock Error Analysis

As mentioned in Section 5.4, we automatically detected and manually fixed a number of mistakes that we found in the original chord annotations. In some instances, although the root of the provided chord label was missing from the corresponding segment, the label was in fact correct. In these instances, it was often the case that the root appeared in the previous segment and thus was still perceptually salient to the listener, either because of its long duration or because it appeared in the last event of the previous segment. Sometimes, the same harmonic and melodic patterns were repeated throughout the piece, with the root appearing in the first few repetitions of these patterns,




Figure 6: Measures 14-15 of ‘Let It Be’ by the Beatles, where HMPerception incorrectly predicts G:maj6 for measure 15 (bottom), while semi-CRF correctly predicts G:maj (top).

but disappearing later on. This was true for ‘Twist and Shout’ by the Beatles, in which the same I IV V7 progression of C major, F major, and G dominant 7 is repeated throughout the song, with the root C disappearing from C major segments by measure 11. Due to their inability to exploit larger scale patterns, neither system could predict the correct label for such segments.

We also found that three of the songs that we manually detected as having labels with incorrect modes (‘Great Balls of Fire,’ ‘Heartbreak Hotel,’ and ‘Shake, Rattle, and Roll’) were heavily influenced by blues. The three songs contain many major chord segments where the major third is purposefully swapped for a minor third to create a blues feel. We kept the labels as they were in these instances, but again both systems struggled to correctly predict the true label in these cases.

Figure 6 contains a brief excerpt from ‘Let It Be’ by the Beatles demonstrating the utility of a segmental approach over an event-based approach. Semi-CRF correctly predicts a segment spanning measure 15 with the label G:maj, while HMPerception predicts these same segment boundaries, but incorrectly produces the label G:maj:add6. Semi-CRF most likely predicts the correct label because of its ability to heuristically detect figuration: the E5 on the first beat of measure 15 is a suspension, while the E5 on the fourth beat is a neighboring tone. It would be difficult for an event-based approach to recognize these notes as nonharmonic tones, as detecting figuration requires segment information. For instance, to detect a neighbor, this requires determining if one of its anchor notes belongs to the candidate segment (see Appendix B for a full definition of neighbor and anchor tones).

7. Related Work

Numerous approaches for computerized harmonic analysis have been proposed over the years, starting with the pioneering system of Winograd (1968), in which a systemic grammar was used to encode knowledge of harmony. Barthelemy and Bonardi (2001) and more recently Rizo et al. (2016) provide a good survey of previous work in harmonic analysis of symbolic music. Here, we focus on the three systems that

inspired our work: Melisma (Temperley and Sleator, 1999), HarmAn (Pardo and Birmingham, 2002), and HMPerception (Radicioni and Esposito, 2010) (listed in chronological order). These systems, as well as our semi-CRF approach, incorporate knowledge of music theory through manually defined *rules* or *features*. For example, the “compatibility rule” used in Melisma is analogous to the chord coverage features used in the semi-CRF, the “positive evidence” score computed based on the six template classes in HarmAn, or the “Asserted-notes” features in HMPerception. Likewise, the segment purity features used in semi-CRF are analogous to the “negative evidence” scores from HarmAn, while the figuration heuristics used in semi-CRF can be seen as the counterpart of the “ornamental dissonance rule” used in Melisma. In these systems, each rule or feature is assigned an importance, or weight, in order to enable the calculation of an overall score for any candidate chord segmentation. Given a set of weights, optimization algorithms are used to determine the maximum scoring segmentation and labeling of the musical input. HMPerception uses the Viterbi algorithm (Rabiner, 1989) to find the optimal sequence of event labels, whereas semi-CRF uses a generalization of Viterbi (Sarawagi and Cohen, 2004) to find the joint most likely segmentation and labeling. The dynamic programming algorithm used in Melisma is actually an instantiation of the same general Viterbi algorithm – like HMPerception and semi-CRF it makes a first-order Markov assumption and computes a similar lattice structure that enables a linear time complexity in the length of the input. HarmAn, on the other hand, uses the Relaxation algorithm (Cormen et al., 2009), whose original quadratic complexity is reduced to linear through a greedy approximation.

While the four systems are similar in terms of the musical knowledge they incorporate and their optimization algorithms, there are two important aspects that differentiate them:

1. Are the weights learned from the data, or pre-specified by an expert? HMPerception and semi-CRF train their parameters, whereas Melisma and HarmAn have parameters that are predefined manually.
2. Is chord recognition done as a joint segmentation and labeling of the input, or as a labeling of event sequences? HarmAn and semi-CRF are in the segment-based labeling category, whereas Melisma and HMPerception are event-based.

Learning the weights from the data is more feasible, more scalable, and, given a sufficient amount of training examples, much more likely to lead to optimal performance. Furthermore, the segment-level classification has the advantage of enabling segment-level features that can be more informative than event-level analogues. The semi-CRF approach described in this paper is the first to take advantage of both learning

the weights and performing a joint segmentation and labeling of the input.

8. Future Work

Manually engineering features for chord recognition is a cognitively demanding and time consuming process that requires music theoretical knowledge and that is not guaranteed to lead to optimal performance, especially when complex features are required. In future work we plan to investigate automatic feature extraction using recurrent neural networks (RNN). While RNNs can theoretically learn useful features from raw musical input, they are still event-level taggers, even when used in more sophisticated configurations, such as bi-directional deep LSTMs (Graves, 2012). We plan to use the Segmental RNNs of Kong et al. (2016), which combine the benefits of RNNs and semi-CRFs: bidirectional RNNs compute representations of candidate segments, whereas segment-label compatibility scores are integrated using a semi-Markov CRF. Learning the features entirely from scratch could require a larger number of training examples, which may not be feasible to obtain. An alternative is to combine RNN sequence models with explicit knowledge of music theory, as was done recently by Jaques et al. (2017) for the task of melody generation.

Music analysis tasks are mutually dependent on each other. Voice separation and chord recognition, for example, have interdependencies, such as figuration notes belonging to the same voice as their anchor notes. Temperley and Sleator (1999) note that harmonic analysis, in particular chord changes, can benefit meter modeling, whereas knowledge of meter is deemed crucial for chord recognition. This “serious chicken-and-egg problem” can be addressed by modeling the interdependent tasks together, for which probabilistic graphical models are a natural choice. Correspondingly, we plan to develop models that jointly solve multiple music analysis tasks, an approach that reflects more closely the way humans process music.

9. Conclusion

We presented a semi-Markov CRF model that approaches chord recognition as a joint segmentation and labeling task. Compared to event-level tagging approaches based on HMMs or linear CRFs, the segment-level approach has the advantage that it can accommodate features that consider all the notes in a candidate segment. This capability was shown to be especially useful for music with complex textures that diverge from the simpler note-for-note structures of the Bach chorales. The semi-CRF’s parameters are trained on music annotated with chord labels, a data-driven approach that is more feasible than manually tuning the parameters, especially when the number of rules or features is large. Empirical evaluations on three datasets of classical music and a newly created

dataset of rock music show that the semi-CRF model performs substantially better than previous approaches when trained on a sufficient number of labeled examples and stays competitive when the training data is small. The code is made publicly available on the first author’s GitHub⁷.

Notes

¹ Link to TAVERN:

<https://github.com/jcdevaney/TAVERN>

² Link to Humdrum:

<http://www.humdrum.org/Humdrum/representations/harm.rep.html>

³ Link to Kostka-Payne corpus:

<http://www.cs.northwestern.edu/~pardo/kpcorpus.zip>

⁴ Link to PhotoScore:

<http://www.neuratron.com/photoscore.htm>

⁵ Link to StatNLP:

<http://statnlp.org/research/ie/>

⁶ Link to David Temperley’s Melisma Music Analyzer:

<http://www.link.cs.cmu.edu/melisma/>

⁷ Link to Code:

https://github.com/kristenmasada/chord_recognition_semi_crf

Acknowledgments

We would like to thank Patrick Gray for his help with pre-processing the TAVERN corpus. We thank Daniele P. Radicioni and Roberto Esposito for their help and their willingness to share the original BaCh dataset and the HMPerception implementation. We would like to thank Bryan Pardo for sharing the KP Corpus and David Temperley for providing us with an updated version that fixed some of the labeling errors in the original KP corpus. We would also like to thank David Chelberg for providing valuable comments on a previous version of the manuscript. Finally, we would like to thank the anonymous reviewers and the editorial team for their insightful comments.

Kristen Masada’s work on this project was partly supported by an undergraduate research apprenticeship grant from the Honors Tutorial College at Ohio University.

References

- Aldwell, E., Schachter, C., and Cadwallader, A. (2011). *Harmony and Voice Leading*. Schirmer, 4th edition.
- Baeza-Yates, R. A. and Ribeiro-Neto, B. (1999). *Modern Information Retrieval*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA.
- Barthelemy, J. and Bonardi, A. (2001). Figured bass and tonality recognition. In *International Symposium on Music Information Retrieval (ISMIR)*.
- Collins, M. (2002). Discriminative training methods for hidden markov models: Theory and experi-

- ments with perceptron algorithms. In *Proceedings of the ACL-02 Conference on Empirical Methods in Natural Language Processing*, EMNLP '02, pages 1–8, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Cormen, T. H., Leiserson, C. E., Rivest, R. L., and Stein, C. (2009). *Introduction to Algorithms, Third Edition*. The MIT Press, 3rd edition.
- Denyer, R. (1992). *The Guitar Handbook: A Unique Source Book for the Guitar Player*. Knopf.
- Devaney, J., Arthur, C., Condit-Schultz, N., and Nisula, K. (2015). Theme and variation encodings with roman numerals (TAVERN): A new data set for symbolic music analysis. In *International Society for Music Information Retrieval Conference (ISMIR)*.
- Gómez, E. (2006). *Tonal Description of Music Audio Signals*. PhD thesis, Universitat Pompeu Fabra.
- Graves, A. (2012). *Supervised Sequence Labelling with Recurrent Neural Networks*, volume 385 of *Studies in Computational Intelligence*. Springer.
- Harte, C. (2010). *Towards Automatic Extraction of Harmony Information from Music Signals*. PhD thesis, Queen Mary University of London.
- Jaques, N., Gu, S., Bahdanau, D., Hernández-Lobato, J. M., Turner, R. E., and Eck, D. (2017). Sequence tutor: Conservative fine-tuning of sequence generation models with kl-control. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, pages 1645–1654, Sydney, NSW, Australia.
- Kong, L., Dyer, C., and Smith, N. A. (2016). Segmental recurrent neural networks. In *International Conference on Learning Representations (ICLR)*, San Juan, Puerto Rico.
- Kostka, S. and Payne, D. (1984). *Tonal Harmony*. McGraw-Hill.
- Kschischang, F. R., Frey, B., and Loeliger, H.-A. (2001). Factor graphs and the sum-product algorithm. *IEEE Transactions on Information Theory*, 47(2):498–519.
- Lafferty, J., McCallum, A., and Pereira, F. (2001). Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Proceedings of the Eighteenth International Conference on Machine Learning*, pages 282–289, Williamstown, MA.
- Maxwell, H. J. (1992). An expert system for harmonizing analysis of tonal music. In Balaban, M., Ebcioğlu, K., and Laske, O., editors, *Understanding Music with AI*, pages 334–353. MIT Press, Cambridge, MA, USA.
- Muis, A. O. and Lu, W. (2016). Weak semi-Markov CRFs for noun phrase chunking in informal text. In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 714–719, San Diego, California. Association for Computational Linguistics.
- Papadopoulos, H. and Peeters, G. (2009). Local key estimation based on harmonic and metric structures. In *International Conference on Digital Audio Effects (DAFx-09)*.
- Pardo, B. and Birmingham, W. P. (2002). Algorithms for chordal analysis. *Computer Music Journal*, 26(2):27–49.
- Rabiner, L. R. (1989). A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286.
- Radicioni, D. P. and Esposito, R. (2010). BREVE: an HMPerceptron-based chord recognition system. In *Advances in Music Information Retrieval*, pages 143–164. Springer Berlin Heidelberg.
- Raphael, C. and Stoddard, J. (2003). Harmonic analysis with probabilistic graphical models. In *International Society for Music Information Retrieval Conference (ISMIR)*.
- Rizo, D., Illescas, P. R., and Iñesta, J. M. (2016). Interactive melodic analysis. In Meredith, D., editor, *Computational Music Analysis*, pages 191–219. Springer International Publishing, Cham.
- Rocher, T., Robine, M., Hanna, P., and Strandh, R. (2009). Dynamic chord analysis for symbolic music. In *International Computer Music Conference, ICMC*.
- Sarawagi, S. and Cohen, W. W. (2004). Semi-Markov Conditional Random Fields for information extraction. In *Proceedings of the 17th International Conference on Neural Information Processing Systems, NIPS'04*, pages 1185–1192, Cambridge, MA, USA. MIT Press.
- Scholz, R. and Ramalho, G. (2008). Cochonut: Recognizing complex chords from midi guitar sequences. In *International Conference on Music Information Retrieval (ISMIR)*, pages 27–32, Philadelphia, USA.
- Taylor, E. (1989). *The AB Guide to Music Theory Part 1*. ABRSM.
- Temperley, D. and Sleator, D. (1999). Modeling meter and harmony: A preference-rule approach. *Computer Music Journal*, 23(1):10–27.
- Tsochantaridis, I., Hofmann, T., Joachims, T., and Altun, Y. (2004). Support vector machine learning for interdependent and structured output spaces. In *Proceedings of the Twenty-first International Conference on Machine Learning, ICML '04*, pages 104–111, New York, NY, USA. ACM.
- Winograd, T. (1968). Linguistics and the computer analysis of tonal harmony. *Journal of Music Theory*, 12(1):2–49.

A. Types of Chords in Tonal Music

A chord is a group of notes that form a cohesive harmonic unit to the listener when sounding simulta-

neously (Aldwell et al., 2011). We design our system to handle the following types of chords: triads, augmented 6th chords, suspended chords, and power chords.

A.1 Triads

A *triad* is the prototypical instance of a chord. It is based on a root note, which forms the lowest note of a chord in standard position. A third and a fifth are then built on top of this root to create a three-note chord. Inverted triads also exist, where the third or fifth instead appears as the lowest note. The chord labels used in our system do not distinguish among inversions of the same chord. However, once the basic triad is determined by the system, finding its inversion can be done in a straightforward post-processing step, as a function of the bass note in the chord. The quality of the third and fifth intervals of a chord in standard position determines the mode of a triad. For our system, we consider three triad modes: *major* (maj), *minor* (min), and *diminished* (dim). A major triad consists of a major third interval (i.e. 4 half steps) between the root and third, as well as a perfect fifth (7 half steps) between the root and fifth. A minor triad has a minor third interval (3 half steps) between the root and third. Lastly, a diminished triad maintains the minor third between the root and third, but contains a diminished fifth (6 half steps) between the root and fifth. Figure 7 shows these three triad modes, together with three versions of a C major chord, one for each possible type of added note, as explained below.




Figure 7: Triads in 3 modes and with 3 added notes.

A triad can contain an added note, or a fourth note. We include three possible added notes in our system: a fourth, a sixth, and a seventh. A fourth chord (add4) contains an interval of a perfect fourth (5 half steps) between the root and the added note for all modes. In contrast, the interval between the root and added note of a sixth chord (add6) of any mode is a major sixth (9 half steps). For seventh chords (add7), the added note interval varies. If the triad is major, the added note can form a major seventh (11 half steps) with the root, called a major seventh chord. It can also form a minor seventh (10 half steps) to create a dominant seventh chord. If the triad is minor, the added seventh can again either form an interval of a major seventh, creating a minor-major seventh chord, or a minor seventh, forming a minor seventh chord. Finally, diminished triads most frequently contain a diminished seventh interval (9 half steps), producing a fully diminished seventh chord, or a minor seventh interval,

creating a half-diminished seventh chord.

A.2 Augmented 6th Chords

An *augmented 6th chord* is a type of chromatic chord defined by an augmented sixth interval between the lowest and highest notes of the chord (Aldwell et al., 2011). The three most common types of augmented 6th chords are *Italian*, *German*, and *French* sixth chords, as shown in Figure 8 in the key of A minor. In a minor scale, Italian sixth chords can be seen as iv chords with a sharpened root, in the first inversion. Thus, they can be created by stacking the sixth, first, and sharpened fourth scale degrees. In minor, German sixth chords are iv⁷ (i.e. minor seventh) chords with a sharpened root, in the first inversion. They are formed by combining the sixth, first, third, and sharpened fourth scale degrees. Lastly, French sixth chords are created by stacking the sixth, first, second, and sharpened fourth scale degrees. Thus, they are ii^{ø7} (i.e. half-diminished seventh) chords with a sharpened third, in second inversion.



Figure 8: Common types of augmented 6th chords, shown for the A minor scale. The same notes would also be used for the A major scale.

A.3 Suspended and Power Chords

Both *suspended* and *power chords* are similar to triads in that they contain a root and a perfect fifth. They differ, however, in their omission of the third. As shown in Figure 9, *suspended second chords* (sus2) use a second as replacement for this third, forming a major second (2 half steps) with the root, while *suspended fourth chords* (sus4) employ a perfect fourth as replacement (Taylor, 1989). The suspended second and fourth often resolve to a more stable third. In addition to these two kinds of suspended chords, our system considers suspended fourth chords that contain an added minor seventh, forming a *dominant seventh suspended fourth chord* (7sus4).




Figure 9: Suspended and power chords.

In contrast with suspended chords, power chords (pow) do not contain a replacement for the missing third. They simply consist of a root and a perfect fifth. Though they are not formally considered to be chords in classical music, they are commonly referred to in both rock and pop music (Denyer, 1992).

A.4 Chord Ambiguity

Sometimes, the same set of notes can have multiple chord interpretations. For example, the German sixth chord shown in Figure 8 can also be interpreted as an F dominant seventh chord. Added notes can also lead to other types of ambiguity, for example {D, F, A, C} could be an F major sixth chord (i.e. F major with an added sixth) or a D minor seventh chord (i.e. D minor chord with an added minor seventh). Human annotators can determine the correct chord interpretation based on cues such as inversions and context. The semi-CRF model described in this paper captures inversions through the bass features (Appendix C.3), whereas context is taken into account through the chord bigram features (Appendix C.4). This could be further improved by adding other features, such as determining how notes in the current chord resolve to notes in the next chord.

B. Figuration Heuristics

We designed a set of heuristics to determine whether a note n from a segment s is a figuration note with respect to a candidate chord label y . The heuristic rules shown below discover four types of figurations: passing and neighbor notes (Figure 10), and suspensions and anticipations (Figure 11).



Figure 10: Examples of figuration notes, in red.

Passing: There are two anchor notes n_1 and n_2 such that: n_1 's offset coincides with n 's onset; n_2 's onset coincides with n 's offset; n_1 is one scale step below n and n_2 is one step above n , or n_1 is one step above n and n_2 one step below; n is not longer than either n_1 or n_2 ; the accent value of n is strictly smaller than the accent value of n_1 ; at least one of the two anchor notes belongs to segment s ; n is non-harmonic with respect to chord y , i.e. n is not equivalent to the root, third, fifth, or added note of y ; both n_1 and n_2 are harmonic with respect to the segments they belong to.

Neighbor: There are two anchor notes n_1 and n_2 such that: n_1 's offset coincides with n 's onset; n_2 's onset coincides with n 's offset; n_1 and n_2 are both either one step below or one step above n ; n is not longer than either n_1 or n_2 ; the accent value of n is strictly smaller than the accent value of n_1 ; at least one of the two anchor notes belongs to segment s ; n is non-harmonic with respect to chord y ; both anchor notes are harmonic with respect to the segments they belong to.

Suspension: Note n belongs to the first event of segment s . There is an anchor note m in the previous event (last event in the previous segment) such that: m



Figure 11: Examples of figuration notes, in red.

and n have the same pitch; n is either tied with m (i.e. held over) or m 's offset coincides with n 's onset (i.e. restruck); n is not longer than m ; n is non-harmonic with respect to chord y , while m is harmonic with respect to the previous chord.

Anticipation: Note n belongs to the last event of segment s . There is an anchor note m in the next event (first event in the next segment) such that: n and m have the same pitch; m is either tied with n (i.e. held over) or n 's offset coincides with m 's onset (i.e. restruck); n is not longer than m ; n is non-harmonic with respect to chord y , while m is harmonic relative to all other notes in its event.

Furthermore, because the weak semi-CRF features shown in Equation 4 do not have access to the candidate label y_{k-1} of the previous segment s_{k-1} , we need a heuristic to determine whether an anchor note is harmonic whenever the anchor note belongs to the previous segment. The heuristic simply looks at the other notes in the event containing the anchor note: if the event contains 2 or more other notes, at least 2 of them need to be consonant with the anchor, i.e. intervals of octaves, fifths, thirds, and their inversions; if the event contains just one note other than the anchor note, it has to be consonant with the anchor.

We emphasize that the rules mentioned above for detecting figuration notes are only approximations. We recognize that correctly identifying figuration notes can also depend on subtler stylistic and contextual cues, thus allowing for exceptions to each of these rules.

C. Chord Recognition Features

Given a segment s and chord y , we will use the following notation:

- $s.Notes$, $s.N$ = the set of notes in the segment s .
- $s.Events$, $s.E$ = the sequence of events in s .
- $e.len$, $n.len$ = the length (i.e. duration) of event e or note n , in quarters.
- $e.acc$, $n.acc$ = the accent value of event e or note n , as computed by the `beatStrength()` function in Music21².
- $y.root$, $y.third$, and $y.fifth$ = the triad tones of the chord y .
- $y.added$ = the added note of chord y , if y is an added tone chord.
- $s.Fig(y)$ = the set of notes in s that are *figuration* with respect to chord y .

²Link to Music21: <http://web.mit.edu/music21>

- $s.NonFig(y) = s.Notes - s.Fig(y)$ = the set of notes in s that are not figuration with respect to y .

Note that a note may contain multiple events, as such the note length $n.len$ can be seen as the sum of the length of all events that span the duration of that note. For example, the first G3 in the bass of Figure 2 has a length of a quarter – it corresponds to the G3 in measure 2 of Figure 1 and is shown as a tied note to simplify the description. Therefore its $n.len = 1$. Each of the two events that span its duration have a length of an eighth, hence $e_1.len = e_2.len = 0.5$.

The *accent value* is determined based on the metrical position of a note or event, e.g. in a song written in a 4/4 time signature, the first beat position would have a value of 1.0, the third beat 0.5, and the second and fourth beats 0.25. Any other eighth note position within a beat would have a value of 0.125, any sixteenth note position strictly within the beat would have a value of 0.0625, and so on. To determine whether a note n from a segment s is a *figuration* note with respect to a candidate chord label y , we use a set of heuristics, as detailed in Appendix B.

The duration and accent-weighted segment-level features introduced in this section have real values. Given a real-valued feature $f(s, y)$ that takes values in $[0, 1]$, we discretize it into $K+2$ Boolean features by partitioning the $[0, 1]$ interval into a set of K subinterval bins $\mathcal{B} = \{(b_{k-1}, b_k] | 1 \leq k \leq K\}$. For each bin, the corresponding Boolean feature determines whether $f(s, y) \in (b_{k-1}, b_k]$. Additionally, two Boolean features are defined for the boundary cases $f(s, y) = 0$ and $f(s, y) = 1$. For each real-valued feature, unless specified otherwise, we use the bin set $\mathcal{B} = [0, 0.1, \dots, 0.9, 1.0]$.

C.1 Segment Purity

The segment purity feature $f_1(s, y)$ computes the fraction of the notes in segment s that are harmonic, i.e. belong to chord y :

$$f_1(s, y) = \frac{\sum_{n \in s.Notes} \mathbf{1}[n \in y]}{|s.Notes|}$$

The duration-weighted version $f_2(s, y)$ of the purity feature weighs each note n by its length $n.len$:

$$f_2(s, y) = \frac{\sum_{n \in s.Notes} \mathbf{1}[n \in y] * n.len}{\sum_{n \in s.Notes} n.len}$$

The accent-weighted version $f_3(s, y)$ of the purity feature weighs each note n by its accent weight $n.acc$:

$$f_3(s, y) = \frac{\sum_{n \in s.Notes} \mathbf{1}[n \in y] * n.acc}{\sum_{n \in s.Notes} n.acc}$$

The 3 real-valued features are discretized using the default bin set \mathcal{B} .

C.1.1 Figuration-Controlled Segment Purity

For each segment purity feature, we create a figuration-controlled version that ignores notes that were heuristically detected as figuration, i.e. replace $s.Notes$ with $s.NonFig(y)$ in each feature formula.

C.2 Chord Coverage

The chord coverage features determine which of the chord notes belong to the segment. In this section, each of the coverage features are non-zero only for major, minor, and diminished triads and their added note counterparts. This is implemented by first defining an indicator function $y.Triad$ that is 1 only for triads and chords with added notes, and then multiplying it into all the triad features from this section.

$$y.Triad = \mathbf{1}[y.mode \in \{\text{maj}, \text{min}, \text{dim}\}]$$

Furthermore, we compress notation by showing the mode predicates as attributes of the label, e.g. $y.maj$ is a predicate equivalent with testing whether $y.mode = \text{maj}$. Thus, an equivalent formulation of $y.Triad$ is as follows:

$$y.Triad = \mathbf{1}[y.maj \vee y.min \vee y.dim]$$

To avoid clutter, we do not show $y.Triad$ in any of the features below, although it is assumed to be multiplied into all of them. The first 3 coverage features refer to the triad notes:

$$\begin{aligned} f_4(s, y) &= \mathbf{1}[y.root \in s.Notes] \\ f_5(s, y) &= \mathbf{1}[y.third \in s.Notes] \\ f_6(s, y) &= \mathbf{1}[y.fifth \in s.Notes] \end{aligned}$$

A separate feature determines if the segment contains all the notes in the chord:

$$f_7(s, y) = \prod_{n \in y} \mathbf{1}[n \in s.Notes]$$

A chord may have an added tone $y.added$, such as a 4th, a 6th, or a 7th. If a chord has an added tone, we define two features that determine whether the segment contains the added note:

$$\begin{aligned} f_8(s, y) &= \mathbf{1}[\exists y.added \wedge y.added \in s.Notes] \\ f_9(s, y) &= \mathbf{1}[\exists y.added \wedge y.added \notin s.Notes] \end{aligned}$$

Through the first feature, the system can learn to prefer the added tone version of the chord when the segment contains it, while the second feature enables the system to learn to prefer the triad-only version if no added tone is in the segment. To prevent the system from recognizing added chords too liberally, we add a feature that is triggered whenever the total length of the added notes in the segment is greater than the total length of the root:

$$alen(s, y) = \sum_{n \in s.Notes} \mathbf{1}[n = y.added] * n.len$$

$$rlen(s, y) = \sum_{n \in s.Notes} \mathbf{1}[n = y.root] * n.len$$

$$f_{10}(s, y) = \mathbf{1}[\exists y.added] * \mathbf{1}[alen(s, y) > rlen(s, y)]$$

The duration-weighted versions of the chord coverage features weigh each chord tone by its total duration in the segment. For the root, the feature would be computed as shown below:

$$f_{11}(s, y) = \frac{\sum_{n \in s.Notes} \mathbf{1}[n = y.root] * n.len}{\sum_{n \in s.Notes} n.len}$$

Similar features f_{12} and f_{13} are computed for the third and the fifth. The corresponding accent-weighted features f_{14} , f_{15} , and f_{16} are computed in a similar way, by replacing the note duration $n.len$ in the duration-weighted formulas with the note accent value $n.acc$.

The duration-weighted feature for the added tone is computed similarly:

$$f_{17}(s, y) = \frac{\mathbf{1}[\exists y.added] * \sum_{n \in s.Notes} \mathbf{1}[n = y.added] * n.len}{\sum_{n \in s.Notes} n.len}$$

Furthermore, by replacing $n.len$ with $n.acc$, we also obtain the accent-weighted version f_{18} .

An alternative definition of duration-weighted features is based on the proportion of the segment time that is covered by a particular chord note. The corresponding duration-weighted feature for the chord root is shown below:

$$f_{19}(s, y) = \frac{\sum_{e \in s.Events} \mathbf{1}[y.root \in e] * e.len}{\sum_{e \in s.Events} e.len}$$

Similar duration-weighted features normalized by the segment length are defined for thirds, fifths, and added notes.

All duration-weighted and accent-weighted features are discretized using the default bin set \mathcal{B} .

C.2.1 Chord Coverage for Augmented 6th Chords

We label each note appearing in an augmented 6th chord as follows:

- $y.bass$ = the lowest note.
- $y.3rd$ = the note that is a third above $y.bass$.
- $y.6th$ = the note that is an augmented sixth above $y.bass$.
- $y.5th$ = the note that forms a perfect fifth (for German 6th chords) or a diminished fifth (for French) above $y.bass$.

The features defined in this section are non-zero only for augmented 6th chord labels. Similar to Appendix C.2, we define an indicator function $y.AS$ that is 1 only for augmented 6th chords and implicitly multiply this into each of the features from this section.

$$\begin{aligned} y.AS &= \mathbf{1}[y.mode \in \{it6, fr6, ger6\}] \\ y.AS &= \mathbf{1}[y.it6 \vee y.fr6 \vee y.ger6] \end{aligned}$$

We define an additional indicator function $y.FG$ that is 1 only for French and German 6th chords.

$$y.FG = \mathbf{1}[y.fr6 \vee y.ger6]$$

The coverage features for augmented 6th chords are overall analogous to the ones for triad chords.

$$\begin{aligned} as_1(s, y) &= \mathbf{1}[y.bass \in s.Notes] \\ as_2(s, y) &= \mathbf{1}[y.3rd \in s.Notes] \\ as_3(s, y) &= \mathbf{1}[y.6th \in s.Notes] \\ as_4(s, y) &= \mathbf{1}[y.FG \wedge y.5th \in s.Notes] \end{aligned}$$

The duration-weighted versions are as follows:

$$as_5(s, y) = \frac{\sum_{n \in s.Notes} \mathbf{1}[n = y.bass] * n.len}{\sum_{n \in s.Notes} n.len}$$

$$as_6(s, y) = \mathbf{1}[y.FG] * \frac{\sum_{n \in s.Notes} \mathbf{1}[n = y.5th] * n.len}{\sum_{n \in s.Notes} n.len}$$

As before, we replace $n.len$ with $n.acc$ to obtain the accent-weighted versions of as_5 and as_6 . We also define segment-based duration-weighted features:

$$as_7(s, y) = \frac{\sum_{e \in s.Events} \mathbf{1}[y.bass \in e] * e.len}{\sum_{e \in s.Events} e.len}$$

C.2.2 Chord Coverage for Suspended and Power Chords

As before, we define the features in this section to be non-zero only for suspended or power chord labels by implicitly multiplying them with an indicator function $y.SP$.

$$y.SP = \mathbf{1}[y.sus2 \vee y.sus4 \vee y.7sus4 \vee y.pow]$$

The coverage features for suspended and power chords are also similar to the ones defined for triad chords.

$$\begin{aligned} sp_1(s, y) &= \mathbf{1}[y.root \in s.Notes] \\ sp_2(s, y) &= \mathbf{1}[y.sus2 \wedge y.2nd \in s.Notes \vee \\ &\quad (y.sus4 \vee y.7sus4) \wedge y.4th \in s.Notes] \\ sp_3(s, y) &= \mathbf{1}[y.5th \in s.Notes] \\ sp_4(s, y) &= \mathbf{1}[y.7sus4 \wedge y.7th \in s.Notes] \\ sp_5(s, y) &= \mathbf{1}[y.7sus4 \wedge y.7th \notin s.Notes] \end{aligned}$$

$$alen(s, y) = \sum_{n \in s.Notes} \mathbf{1}[n = y.7th] * n.len$$

$$rlen(s, y) = \sum_{n \in s.Notes} \mathbf{1}[n = y.root] * n.len$$

$$sp_6(s, y) = \mathbf{1}[y.7sus4 \wedge alen(s, y) > rlen(s, y)]$$

The duration-weighted versions are as follows:

$$sp_7(s, y) = \frac{rlen(s, y)}{\sum_{n \in s.Notes} n.len}$$

$$sp_8(s, y) = \mathbf{1}[y.7sus4] * \frac{alen(s, y)}{\sum_{n \in s.Notes} n.len}$$

We also define accent-weighted versions of sp_7 and sp_8 , as well as segment-based duration-weighted features:

$$sp_9(s, y) = \frac{\sum_{e \in s.Events} \mathbf{1}[y.root \in e] * e.len}{\sum_{e \in s.Events} e.len}$$

C.2.3 Figuration-Controlled Chord Coverage

For each chord coverage feature, we create a figuration-controlled version that ignores notes that were heuristically detected as figuration, i.e. replace $s.Notes$ with $s.NonFig(y)$ in each feature formula.

C.3 Bass

The bass note provides the foundation for the harmony of a musical segment. For a correct segment, its bass note often matches the root of its chord label. If the bass note instead matches the chord's third or fifth, or is an added dissonance, this may indicate that the chord is inverted. Thus, comparing the bass note with the chord tones can provide useful features for determining whether a segment is compatible with a chord label. As in Appendix C.2, we implicitly multiply each of these features with $y.Triad$ so that they are non-zero only for triads and chords with added notes.

There are multiple ways to define the bass note of a segment s . One possible definition is the lowest note of the first event in the segment, i.e. $s.e_1.bass$. Comparing it with the root, third, fifth, and added tones of a chord results in the following features:

$$f_{20}(s, y) = \mathbf{1}[s.e_1.bass = y.root]$$

$$f_{21}(s, y) = \mathbf{1}[s.e_1.bass = y.third]$$

$$f_{22}(s, y) = \mathbf{1}[s.e_1.bass = y.fifth]$$

$$f_{23}(s, y) = \mathbf{1}[\exists y.added \wedge s.e_1.bass = y.added]$$

An alternative definition of the bass note of a segment is the lowest note in the entire segment, i.e. $\min_{e \in s.E} e.bass$. The corresponding features will be:

$$f_{24}(s, y) = \mathbf{1}[y.root = \min_{e \in s.E} e.bass]$$

$$f_{25}(s, y) = \mathbf{1}[y.third = \min_{e \in s.E} e.bass]$$

$$f_{26}(s, y) = \mathbf{1}[y.fifth = \min_{e \in s.E} e.bass]$$

$$f_{27}(s, y) = \mathbf{1}[\exists y.added \wedge y.added = \min_{e \in s.E} e.bass]$$

The duration-weighted version of the bass features weigh each chord tone by the time it is used as the

lowest note in each segment event, normalized by the duration of the bass notes in all the events. For the root, the feature is computed as shown below:

$$f_{28}(s, y) = \frac{\sum_{e \in s.Events} \mathbf{1}[e.bass = y.root] * e.len}{\sum_{e \in s.Events} e.len}$$

Similar features f_{29} and f_{30} are computed for the third and the fifth. The duration-weighted feature for the added tone is computed as follows:

$$f_{31}(s, y) = \frac{\mathbf{1}[\exists y.added] * \sum_{e \in s.E} \mathbf{1}[e.bass = y.added] * e.len}{\sum_{e \in s.E} e.len}$$

The corresponding accent-weighted features f_{32} , f_{33} , f_{34} , and f_{35} are computed in a similar way, by replacing the bass duration $e.bass.len$ in the duration-weighted formulas with the note accent value $e.bass.acc$.

All duration-weighted and accent-weighted features are discretized using the default bin set \mathcal{B} .

C.3.1 Bass Features for Augmented 6th Chords

Similar to the chord coverage features in Appendix C.2.1, we assume that the indicator $y.AS$ is multiplied into all features in this section, which means they are non-zero only for augmented 6th chords.

$$as_8(s, y) = \mathbf{1}[s.e_1.bass = y.bass]$$

$$as_9(s, y) = \mathbf{1}[s.e_1.bass = y.3rd]$$

$$as_{10}(s, y) = \mathbf{1}[s.e_1.bass = y.6th]$$

$$as_{11}(s, y) = \mathbf{1}[(y.f r 6 \vee y.ger 6) \wedge s.e_1.bass = y.5th]$$

$$as_{12}(s, y) = \mathbf{1}[y.bass = \min_{e \in s.E} e.bass]$$

$$as_{13}(s, y) = \mathbf{1}[y.3rd = \min_{e \in s.E} e.bass]$$

$$as_{14}(s, y) = \mathbf{1}[y.6th = \min_{e \in s.E} e.bass]$$

$$as_{15}(s, y) = \mathbf{1}[y.FG \wedge y.5th = \min_{e \in s.E} e.bass]$$

We define the following duration-weighted version for the augmented sixth bass and fifth.

$$as_{16}(s, y) = \frac{\sum_{e \in s.E} \mathbf{1}[e.bass = y.bass] * e.len}{\sum_{e \in s.E} e.len}$$

$$as_{17}(s, y) = \mathbf{1}[y.FG] * \frac{\sum_{e \in s.E} \mathbf{1}[e.bass = y.5th] * e.len}{\sum_{e \in s.E} e.len}$$

C.3.2 Bass Features for Suspended and Power Chords

The indicator $y.SP$ is multiplied into all features in this section like in Appendix C.2.2, meaning they are non-

zero only for suspended and power chords.

$$\begin{aligned}
sp_{10}(s, y) &= \mathbf{1}[s.e_1.bass = y.root] \\
sp_{11}(s, y) &= \mathbf{1}[y.sus2 \wedge s.e_1.bass = y.2nd \vee \\
&\quad (y.sus4 \vee y.7sus4) \wedge s.e_1.bass = y.4th] \\
sp_{12}(s, y) &= \mathbf{1}[s.e_1.bass = y.5th] \\
sp_{13}(s, y) &= \mathbf{1}[y.7sus4 \wedge s.e_1.bass = y.7th] \\
\\
sp_{14}(s, y) &= \mathbf{1}[y.root = \min_{e \in s.E} e.bass] \\
sp_{15}(s, y) &= \mathbf{1}[y.sus2 \wedge y.2nd = \min_{e \in s.E} e.bass \vee \\
&\quad (y.sus4 \vee y.7sus4) \wedge y.4th = \min_{e \in s.E} e.bass] \\
sp_{16}(s, y) &= \mathbf{1}[y.5th = \min_{e \in s.E} e.bass] \\
sp_{17}(s, y) &= \mathbf{1}[y.7sus4 \wedge y.7th = \min_{e \in s.E} e.bass]
\end{aligned}$$

The duration-weighted version for the root and seventh are computed as follows:

$$\begin{aligned}
sp_{18}(s, y) &= \frac{\sum_{e \in s.E} \mathbf{1}[e.bass = y.root] * e.len}{\sum_{e \in s.E} e.len} \\
sp_{19}(s, y) &= \mathbf{1}[y.7sus4] * \frac{\sum_{e \in s.E} \mathbf{1}[e.bass = y.7th] * e.len}{\sum_{e \in s.E} e.len}
\end{aligned}$$

C.3.3 Figuration-Controlled Bass

For each bass feature, we create a figuration-controlled version that ignores event bass notes that were heuristically detected as figuration, i.e. replace $e \in s.Events$ with $e \in s.Events \wedge e.bass \notin s.Fig(y)$ in each feature formula.

C.4 Chord Bigrams

The arrangement of chords in chord progressions is an important component of *harmonic syntax* (Aldwell et al., 2011). A first-order semi-Markov CRF model can capture chord sequencing information only through the chord labels y and y' of the current and previous segment. To obtain features that generalize to unseen chord sequences, we follow Radicioni and Esposito (2010) and create chord bigram features using only the *mode*, the *added* note, and the interval in semitones between the roots of the two chords. We define the possible modes of a chord label as follows:

$$\begin{aligned}
\mathcal{M} &= \{\text{maj}, \text{min}, \text{dim}\} \\
&\cup \{\text{it6}, \text{fr6}, \text{ger6}\} \\
&\cup \{\text{sus2}, \text{sus4}, \text{7sus4}, \text{pow}\}
\end{aligned}$$

Other than the common major (maj), minor (min), and diminished (dim) modes, the following chord types have been included in \mathcal{M} as modes:

- Augmented 6th chords: Italian 6th (it6), French 6th (fr6), and German 6th (ger6).

- Suspended chords: suspended second (sus2), suspended fourth (sus4), dominant seventh suspended fourth (7sus4).
- Power (pow) chords.

Correspondingly, the chord bigrams can be generated using the feature template below:

$$\begin{aligned}
g_1(y, y') &= \mathbf{1}[(y.mode, y'.mode) \in \mathcal{M} \times \mathcal{M} \\
&\quad \wedge (y.added, y'.added) \in \{\emptyset, 4, 6, 7\} \times \{\emptyset, 4, 6, 7\} \\
&\quad \wedge |y.root - y'.root| = \{0, 1, \dots, 11\}]
\end{aligned}$$

Note that $y.root$ is replaced with $y.bass$ for augmented 6th chords. Additionally, $y.added$ is always none (\emptyset) for augmented 6th, suspended, and power chords. Thus, $g_1(y, y')$ is a feature template that can generate $(3 \text{ triad modes} \times 4 \text{ added} + 3 \text{ aug6 modes} + 3 \text{ sus modes} + 1 \text{ pow mode})^2 \times 12 \text{ intervals} = 4,332$ distinct features. To reduce the number of features, we use only the $(mode.added)-(mode.added)$ -interval combinations that appear in the manually annotated chord bigrams from the training data.

C.5 Chord Changes and Metrical Accent

In general, repeating a chord creates very little accent, whereas changing a chord tends to attract an accent (Aldwell et al., 2011). Although conflict between meter and harmony is an important compositional resource, in general chord changes support the meter. Correspondingly, a new feature is defined as the accent value of the first event in a candidate segment:

$$f_{36}(s, y) = s.e_1.acc$$